


Speaker ads

The 15-second audio between songs in the store

A shopper hears a person talking, with quiet music underneath. Not a song. Not singing.

Voice ad: two jobs, not a choice




VoxCPM = announcer. FFmpeg = studio mix. Like a Spotify ad, not a song.

- Voice tool talks. Mixer adds licensed music and cuts the file to exactly 15.000 seconds.
- We need both. One tool alone cannot ship the ad.

Order: words, then voice, then mix


Voice pipeline order



VoxCPM does not generate the music bed.

- Only an approved script is turned into voice. That saves money on the rented computer.
- This 15-second file is also the soundtrack of the TV ad. Do not record a second voice.

What is inside the 15 seconds




Spotify-style ad = voice + quiet bed, not singing.

- Voice is loud. Music is quiet. One short sound at the start — not a loop under the whole ad.
- Length is exact so the next song can start on time. If it feels like a song, the music is too loud.

How the app asks for an ad

How the voice API works



One GPU. If /create-video is running, this job waits in queue.

- The app sends the approved script. The rented computer returns one 15-second file.
- If a TV job is using the computer, this voice job waits.

This phase

Voice v1 — what we ship

YES

- ✓ VoxCPM2 Apache-2.0
- ✓ 15s spoken ads
- ✓ Brand voice clone/LoRA later
- ✓ FFmpeg mix with licensed beds
- ✓ POST /create-voice API
- ✓ Queue on one GPU

NO

- ✗ OmniVoice (CC-BY-NC)
- ✗ AI-generated songs
- ✗ ElevenLabs as company-owned core
- ✗ Realtime under 1s required
- ✗ Training and generate on GPU at the same time

If not in YES, do not promise.

Now: spoken 15s ads, licensed music, one rented computer, stop it when idle.

Not now: computer-made songs, singing, or a paid voice vendor as our core.

Full path: ADDLYFT_Full_Pipeline_A_to_Z.pdf